



# Using OpenMP Offload Compilers on Perlmutter GPUs

Helen He, NERSC  
September 4, 2024

# Perlmutter OpenMP Offload Compiler Support

- Vendor provided and supported
  - NVHPC
  - CCE
- Community (Open Source)
  - LLVM/Clang
  - GCC

# Using NVHPC Compiler on Perlmutter GPU

```
% module load PrgEnv-nvidia
```

## **Supports C/C++/Fortran. Use compiler wrappers**

```
# -Minfo is optional which provides compile time info
```

```
% cc -fast -mp=gpu -Minfo=mp,accel src.c
```

```
% CC -fast -mp=gpu -Minfo=mp,accel src.cc
```

```
% ftn -fast -mp=gpu -Minfo=mp,accel src.f90
```

## **Optional runtime messages**

```
% export NVCOMPILER_ACC_NOTIFY=<value>
```

where value is 1: kernel launches 2: data transfers

4: region entry/exit

8: wait operations or synchronizations with the device

16: device memory allocates and deallocates

# Using CCE Compiler on Perlmutter GPU

```
% module load PrgEnv-cray
```

**Supports C/C++. Use compiler wrappers (cc/CC/ftn)**

```
% cc -Ofast -fopenmp src.c
```

```
% CC -Ofast -fopenmp src.cc
```

```
% ftn -O3 -h omp -h noacc src.f90
```

**Optional runtime messages**

```
% export CRAY_ACC_DEBUG=<value>  where value can be 1, 2, 3
```

## Note

- CCE C++ is based on LLVM/Clang; CCE Fortran is classic Cray Compiler
- simd clause needed in CCE Fortran for thread parallelism on GPU

# Using LLVM/Clang Compiler on Perlmutter GPU

```
% module load PrgEnv-llvm
```

## **Supports C/C++. Use native compilers**

```
% clang -Ofast -fopenmp -fopenmp-targets=nvptx64 src.c
```

```
% clang++ -Ofast -fopenmp -fopenmp-targets=nvptx64 src.cc
```

## **Optional runtime messages**

```
% export LIBOMPTARGET_INFO=-1
```

# Using GCC Compiler on Perlmutter GPU

```
% module load PrgEnv-gnu  
% module use /global/cfs/cdirs/nstaff/rgayatri/modules  
% module load gcc-offload/13.2.0
```

## **Supports C/C++/Fortran. Use native compilers**

```
% gcc -Ofast -fopenmp -foffload=nvptx-none="-Ofast -lm -latomic -misa=sm_80" src.c  
% g++ -Ofast -fopenmp -foffload=nvptx-none="-Ofast -lm -latomic -misa=sm_80" src.cc  
% gfortran -Ofast -fopenmp -foffload=nvptx-none="-Ofast -lm -latomic -misa=sm_80" src.f90
```

## **Optional runtime messages**

```
% export GOMP_DEBUG=1
```

## **Warning**

- OpenMP target offload performance not optimal; expect future improvement
- NERSC has done minimal testing only, can not offer good support

# Using Perlmutter GPU

Perlmutter documentation: <https://docs.nersc.gov/systems/perlmutter/>

```
% ssh user-name@perlmutter.nersc.gov
```

```
# Get on a GPU node via salloc:
```

```
% salloc -C gpu -N 1 -q shared -c 32 -G 1 -t 1:00:00 -A ntrain5
```

```
<get on a GPU node>
```

```
% cc -fast -mp=gpu -Minfo=mp,accel mycode.c
```

```
% export OMP_NUM_THREADS=8 # for CPU OpenMP
```

```
./mycode.exe <args>
```

# Sample Batch Script (non-MPI)

```
% cat myjob.sl
#!/bin/bash
#SBATCH -N 1
#SBATCH -C gpu
#SBATCH -t 10:00
#SBATCH -q shared
#SBATCH -c 32
#SBATCH -G 1
#SBATCH -A ntrain5

module load PrgEnv-nvidia
cc -fast -mp=gpu -Minfo=mp,accel -o mycode.exe mycode.c
export OMP_NUM_THREADS=8 # for CPU OpenMP
./mycode.exe <args>
```

```
% module load PrgEnv-nvidia
% sbatch myjob.sl
```

-A is required for GPU jobs  
-q shared -c 32 -G 1 is to share the node



# Perlmutter OpenMP Offload Compilers (as of 9/4/2024)

Compiler	Compile Command	Compile Flags	omp loop	Comment
NVHPC 23.9 C/C++/Fortran	cc/CC/ftn	-fast -mp=gpu	Yes	overall best supported and best performance. <b>Recommend for C/C++ and Fortran codes</b>
CCE 17.0.0 Fortran	ftn	-O3 -h omp -h noacc	Yes	classic Cray Compiler. <b>Recommend for Fortran codes</b>
CCE 17.0.0 C/C++	cc/CC	-Ofast -fopenmp	No	based on LLVM/Clang compiler. <b>Recommend for C/C++ codes</b>
LLVM 18.1.0 C/C++	clang/clang++	-Ofast -fopenmp -fopenmp-targets=nvptx64	No	<b>Recommend for C/C++ codes</b>
GCC 13.2.0 C/C++/Fortran	gcc/g++/gfortran	-Ofast -fopenmp -foffload=nvptx-none="-Ofast -lm -latomic -misa=sm_80"	Yes	performance not optimal as of now