

Session 6: Advanced OpenMP Offload

Christian Terboven, Paul Kapos
IT Center, RWTH Aachen University
Seffenter Weg 23, 52074 Aachen, Germany
{terboven, [kapos](mailto:kapos@itc.rwth-aachen.de)}@itc.rwth-aachen.de

Michael Klemm
AMD; OpenMP ARB
michael.klemm@openmp.org

Yun (Helen) He
NERSC, Lawrence Berkeley National Laboratory, Berkeley, CA, USA
yhe@lbl.gov

Abstract

This document guides you through the exercises. Please follow the instructions given during the training session.

The prepared “Makefile” builds the executables `jacobi.gpu` and `jacobi.sol.gpu` using Nvidia compilers on Perlmutter GPU.

The provided “run.sh” is a batch script that you can use to submit a batch job to run on a Perlmutter GPU compute node. Submit the job via command “`sbatch run.sh`”, and check the output file after it is run, pay attention to total run time.

You can also run an interactive batch job via “`salloc`” to get on a compute node. For example:

```
% salloc -N 1 -q interactive -C gpu -G 1 -c 32 -t 30:00 -A ntrain5 <will land on a compute node>
% module load PrgEnv-nvidia
% make
```

```
% OMP_NUM_THREADS=1 ./jacobi.gpu
% OMP_NUM_THREADS=1 ./jacobi_sol.gpu
```

Please refer to <https://github.com/NERSC/openmp-series-2024/blob/main/Session-5-OpenMP-Offload/Using-OpenMP-Compilers-on-Perlmutter-GPUs-Sept2024.pdf> for more details on Using various OpenMP compilers on Perlmutter GPUs and running jobs.

1. Jacobi on GPU

In this exercise, you will port a Jacobi solver to OpenMP to run on GPUs. This Jacobi example solves a finite difference discretization (5-point-stencil) of the Laplace equation (2D). Please find more info on Jacobi in Session 5 slides deck.

Continue Session 5 tasks 0-3. New task 4 for Session 6.

- Task 0: You might want to acquire reference measurements on the host (i.e., on CPU, without/ GPU)...
- Task 1: Get it to the GPU: Parallelize only the one most compute-intensive loop. Besides using Nvidia compiler, also try to compile the codes using LLVM/Clang compiler, CCE compiler, and GCC compiler.
- Task 2: Improve the data management and the amount of parallelism on the GPU
- Task 3: Optimize that scheduling of iterations for the GPU

Understand the performance of the host and the GPU

- Task 4: Make the code as fast as you can :-).

Use sample codes in exercises/<C, Fortran>/Jacobi2 for hints

Future Tasks: use multiple GPUs, use the host and a GPU, ...